

Backend engineer with 10+ years building and operating the systems that move data — Kafka and RabbitMQ ingest pipelines, Go stream processors, and the high-scale APIs that serve what they land. Comfortable owning a pipeline end to end, from data model and API design through rollout, monitoring, on-call, and incident response, with working depth in Kubernetes-based infrastructure.

EXPERIENCE

Senior Software Engineer

Intuit Mailchimp

Mobile Backend · Landing Pages & Surveys · Email 2

Jun 2024 – Jul 2026 · Atlanta, GA

- Migrated anomaly detection off Data Service Fabric onto AISLE and surfaced industry data on the mobile API root response, modernizing the data path behind the iOS and Android apps.
- Architected and delivered the backend for AI-enhanced audience imports: integrated a preprocessing service into the Imports API, added metadata-based column mapping across the start, column-mapping, and example-data endpoints, and built an asynchronous job runner that fed user corrections back into the model for retraining.
- Modernized the file import pipeline with XLSX support, server-side file-size enforcement on cloud-stored uploads, and a new PATCH endpoint — shipped behind feature flags with fallback to the legacy flow for zero-disruption rollout.
- Expanded Playwright end-to-end coverage across critical account-management flows — signup, forgot username, and forgot password — and fixed IP-throttling middleware that caused intermittent 429s, stabilizing a chronically flaky test pipeline.
- Led the formal root cause analysis for a production incident, then drove the resulting end-to-end testing initiative across the Mobile team, turning a postmortem into a durable change in how the team shipped.
- Took ownership of an under-maintained internal service, adding logging, alerting, and documentation, and authored migration specifications for two services moving onto Intuit's internal service platform — including standing up the team's project, local development environment, and demo resources.
- On Landing Pages & Surveys, retired legacy billing endpoints, upgraded the service stack from AWS SDK v2 to v3, and integrated landing pages into forms listings so published pages were no longer hard to find.
- Drove release management for the next-generation email editor and contributed to the migration parity effort between the legacy and rebuilt editing experiences.

Senior Software Engineer

Roadie

Sender / Geo Team

Sep 2022 – May 2024 · Atlanta, GA

- Built a pipeline ingesting Kafka messages and persisting them to AWS OpenSearch.
- Replaced a legacy tracking API with Benthos, an off-the-shelf Go stream processor, cutting the surface area the team had to maintain.
- Built a location API that proxied and cached Google Maps requests, with fallback logic through HERE when the primary provider failed.
- Contributed to “hasty”, an internal Go package and template that gave the team a fast path to scaffolding new services.
- Triaged outages and handled customer escalations as an on-call engineer.

Senior Software Engineer

Terminus

Data Services Team

Jul 2021 – Aug 2022 · Atlanta, GA

- Implemented features and bug fixes across Go-based data services.
- Migrated legacy visitor data using AWS Athena and Glue.
- Improved monitoring and alerting for data services using Datadog.
- Extended a TypeScript visitor-tracking script to support new event types.

Senior Software Engineer

CNN

Elections · Lego

Dec 2017 – Jul 2021 · Atlanta, GA

- Replaced a legacy Java/Scala platform with a Node.js microservice architecture, and built a RabbitMQ-based messaging system to ingest live voter data.
- Built an AWS serverless API used to register and expose internal services.
- Organized AWS resources with CloudFormation and designed a CI/CD pipeline that simplified deployment and rollback.
- Member of the Elections Technology team that earned an Emmy for Election Night in America.
- Created and maintained internal React and Redux tooling that let producers, operators, and the CNN Political Unit push graphics on air and make real-time calls on election trends.

Software Engineer I–II

CareerBuilder

Job UX & Services · Consumer Intelligence

Jan 2016 – Dec 2017 · Norcross, GA

- Developed and maintained Node.js microservice APIs, persisting large volumes of tracking data across DynamoDB, MySQL, and PostgreSQL.
- Integrated Google's private job search API into CareerBuilder's existing search pages, and built a front-end caching layer for Business Intelligence data.
- Instrumented monitoring and alerting across New Relic, Datadog, and CloudWatch, added a dead letter queue for failed calls in the internal job feeding engine, and built a repeatable deployment process with AWS CodeDeploy.

INFRASTRUCTURE & PROJECTS

Self-Hosted Kubernetes Platform

Personal infrastructure

- Built the platform layer: Gateway API ingress via Envoy, automated TLS through cert-manager and Let's Encrypt DNS-01, DNS records synced to Cloudflare by external-dns, an HA control-plane VIP via kube-vip, a PostgreSQL operator, and NFS-backed persistent storage.
- Operate a 3-node high-availability k3s cluster running roughly 25 self-hosted services, managed entirely through GitOps — every change is a pull request that FluxCD reconciles to cluster state.
- Run the operational stack: Prometheus and Grafana monitoring, uptime checks, SOPS/age-encrypted secrets committed safely to Git, a Tailscale overlay for private access, and an in-cluster Renovate bot that opens dependency update PRs automatically.

SKILLS

Languages	Go, Python, JavaScript, TypeScript, Node.js, PHP, Ruby, SQL, Bash
Backend	REST API design, microservices, Kafka, RabbitMQ, PostgreSQL, MySQL, Redis, DynamoDB, OpenSearch
Infrastructure	Kubernetes (k3s), Docker, FluxCD & GitOps, Helm, Ansible, CloudFormation, AWS, GCP, Envoy, Nginx
Observability	Prometheus, Grafana, Datadog, New Relic, CloudWatch
Frontend	React, Redux, TypeScript, Playwright
Practice	Feature-flagged rollouts, CI/CD, end-to-end test automation, on-call, technical specification authoring

EDUCATION

B.Sc. Computer Science, University of Georgia

2011–2015